

M1.1 Milestone 1.1

```
* Running python /src/m1.1.py
...
EvalMetric: {'accuracy': 0.8673}
0:06.77 elapsed
```

M1.2/1.3 Milestone 1.2/1.3

The GPU spends most of its compute time in

```
* implicit_convolve_sgemm
* sgemm_sm35_ldg_tn_128x8x256x16x32, and
* activation_fw_4d_kernel.
```

Most of the compute time is spend performing convolution. There are also memory API calls that take more time than computation.

```
* Running nvprof python /src/m1.2.py
```

```
...
```

```
==305== Profiling result:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
36.73%	50.515ms	1	50.515ms	50.515ms	50.515ms	implicit_convolve_sgemm
28.66%	39.409ms	1	39.409ms	39.409ms	39.409ms	sgemm_sm35_ldg_tn_128x8x256
14.10%	19.396ms	2	9.6979ms	461.70us	18.934ms	activation_fw_4d_kernel

```
...
```

```
==305== API calls:
```

Time(%)	Time	Calls	Avg	Min	Max	Name
46.82%	1.85071s	18	102.82ms	17.325us	925.01ms	cudaStreamCreateWithFlags
28.48%	1.12580s	10	112.58ms	949ns	320.49ms	cudaFree
20.78%	821.31ms	24	34.221ms	236.10us	814.20ms	cudaMemGetInfo

```
...
```

M2.1 Milestone 2.1

```
* Running python /src/m2.1.py
Loading fashion-mnist data... done
Loading model... done
Op Time: 12.146829
Correctness: 0.8562 Model: ece408-high
```

```
* Running python m2.1.py ece408-low 10000
Loading fashion-mnist data... done
Loading model... done
Op Time: 12.801514
Correctness: 0.629 Model: ece408-low
```

Team Member Contributions

Sujay: Milestone 2 code
Tanishq: Milestone 2 pdf
Gordon: pdf edits

M3.1 Milestone 3.1

```
* Running /usr/bin/time python m3.1.py ece408-high 10000
New Inference
Loading fashion-mnist data... done
Loading model... done
Op Time: 0.509298
```

Correctness: 0.8562 Model: ece408-high
 2.15user 1.05system 0:02.72elapsed 117%CPU (0avgtext+0avgdata 907624maxresident)k
 0inputs+3136outputs (0major+157453minor)pagefaults 0swaps

* Running nvprof python m3.1.py ece408-high 10000

New Inference

Loading fashion-mnist data... done

==309== NVPROF is profiling process 309, command: python m3.1.py ece408-high 10000

Loading model... done

Op Time: 0.556559

Correctness: 0.8562 Model: ece408-high

==309== Profiling application: python m3.1.py ece408-high 10000

==309== Profiling result:

Time(%)	Time	Calls	Avg	Min	Max	Name
83.62%	536.96ms	1	536.96ms	536.96ms	536.96ms	void mxnet::op::forward_kernel<mshadow::gpu, float>(float*, mxnet::op::forward_kernel<mshadow::gpu, float> const *, mxnet::op::forward_kernel<mshadow::gpu, float> const, int, int, int, int, int, int)
6.11%	39.268ms	1	39.268ms	39.268ms	39.268ms	sgemm_sm35_ldg_tn_128x8x256x16x32
3.05%	19.565ms	1	19.565ms	19.565ms	19.565ms	void mshadow::cuda::MapPlanLargeKernel<mshadow::sv::saveto, int=8, int=1024, mshadow::expr::Plan<mshadow::Tensor<mshadow::gpu, int=4, float>, float>, mshadow::expr::Plan<mshadow::expr::BinaryMapExp<mshadow::op::mul, mshadow::expr::ScalarExp<float>, mshadow::Tensor<mshadow::gpu, int=4, float>, float, int=1>, float>>(mshadow::gpu, unsigned int, mshadow::Shape<int=2>, int=4, int)
3.02%	19.393ms	2	9.6963ms	461.59us	18.931ms	void cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>(cudnnTensorStruct, float const *, cudnn::detail::activation_fw_4d_kernel<float, float, int=128, int=1, int=4, cudnn::detail::tanh_func<float>>, cudnnTensorStruct*, float, cudnnTensorStruct*, int, cudnnTensorStruct*)

<snip>

==309== API calls:

Time(%)	Time	Calls	Avg	Min	Max	Name
41.57%	1.86458s	18	103.59ms	17.408us	932.13ms	cudaStreamCreateWithFlags
25.33%	1.13601s	10	113.60ms	1.0270us	325.13ms	cudaFree
18.41%	825.64ms	23	35.898ms	236.27us	818.92ms	cudaMemGetInfo
12.41%	556.53ms	1	556.53ms	556.53ms	556.53ms	cudaDeviceSynchronize

<snip>

Team Member Contributions

Sujay: Milestone 3 pdf

Tanishq: pdf edits

Gordon: Milestone 3 code